




Running Spark on a High-Performance Cluster using RDMA Networking and NVMe Flash


Patrick Stuedi, IBM Research

Hardware Trends

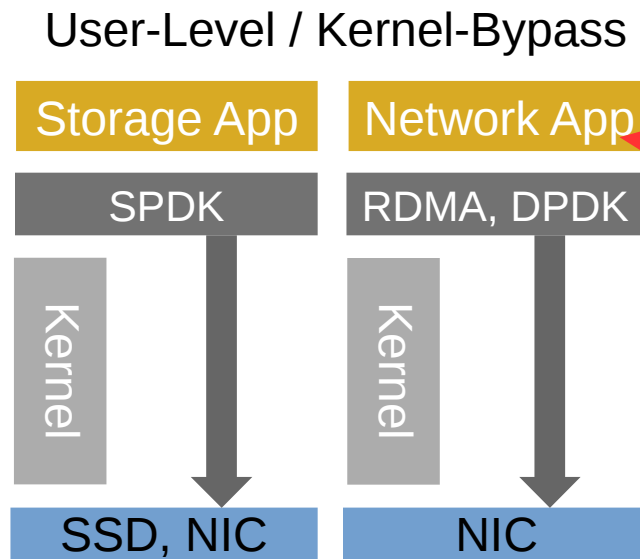

		community target	our target
	2010	2017	2017
Storage	100 MB/s 100ms	1000 MB/s 200us	10 GB/s 50us
Network	1Gbps 50us	10Gbps 20us	100Gbps 2us
CPU	~3GHz	~3GHz	☹



3D XPoint




User-Level APIs




Needed to
achieve
2us RTT!

Remote Data Access


DRAM




NVMe Flash



Let's Use it!

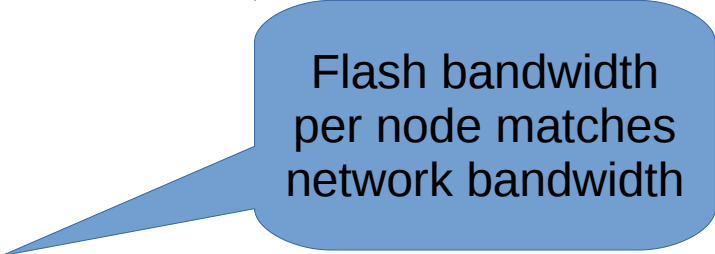


Case Study: Sorting in Spark




Experiment Setup

- Total data size: 12.8 TB
- Cluster size: 128 nodes
- Cluster hardware:
 - DRAM: 512 GB DDR 4
 - Storage: 4x 1.2 TB NVMe SSD
 - Network: 100GbE Mellanox RDMA

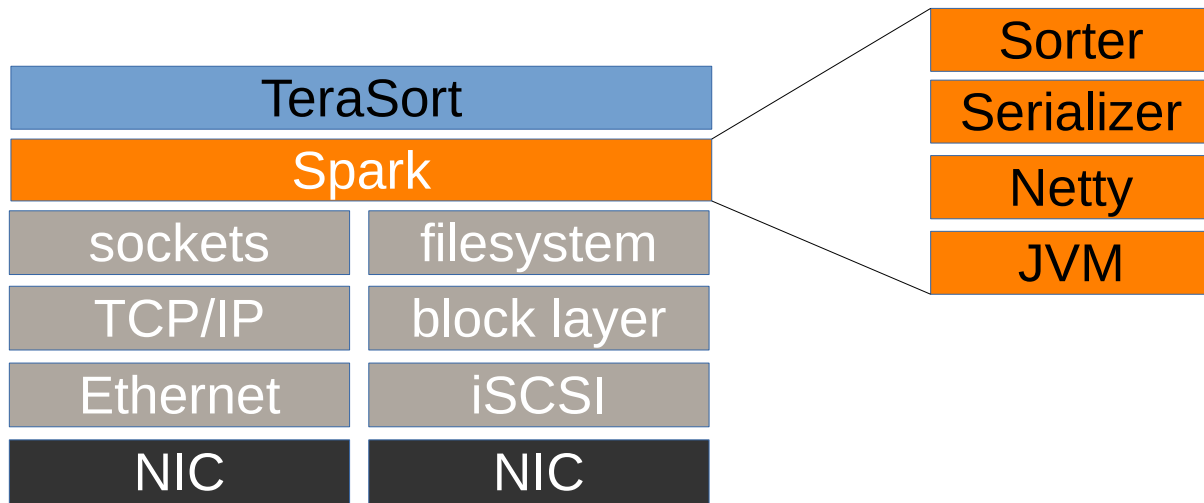


Flash bandwidth
per node matches
network bandwidth

How is the Network Used?




What are the Problems?

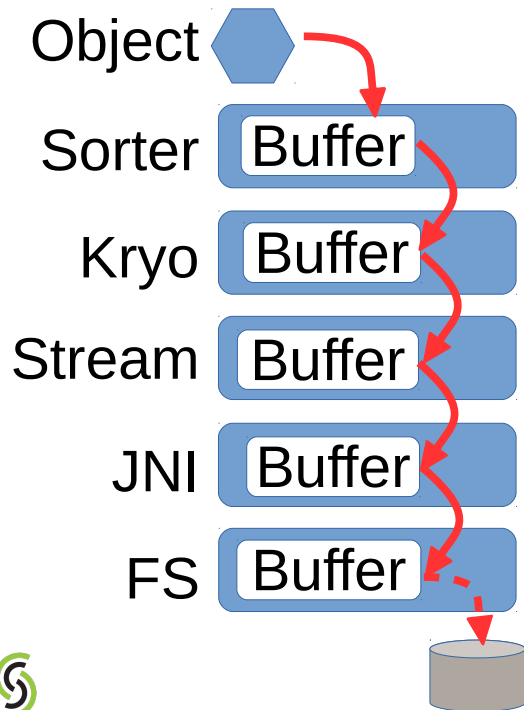


- Spark uses legacy networking and storage APIs: no kernel-bypass
- Spark itself introduces additional I/O layers: Netty, serializer, sorter, etc.


Example: Shuffle (Map)




Example: Shuffle (Map)



Example: Shuffle (Map+Reduce)



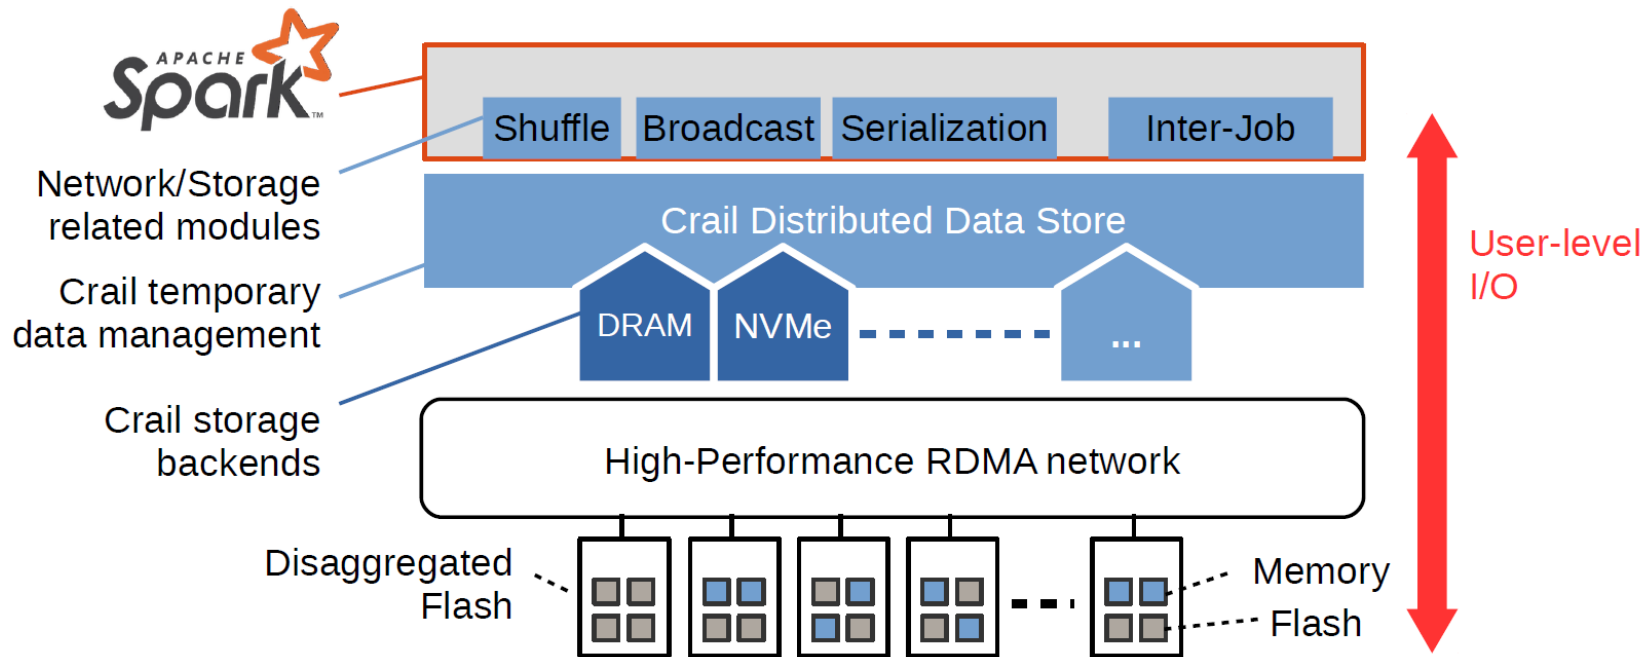
Example: Shuffle (Map+Reduce)



How can we fix this?

- Not just for shuffle
 - Also for broadcast, RDD transport, inter-job sharing, etc.
- Not just for RDMA and NVMe hardware
 - But for any possible future high-performance I/O hardware
- Not just for co-located compute/storage
 - Also for resource disaggregation, heterogeneous resource distribution, etc.
- Not just improve things
 - Make it perform at the hardware limit

The CRAIL Approach





Thank You.

The CRAIL Approach

